Milestone 3 (Week 9): Core Model Development & Integration

Group 5: Aliyyah Jackhan, Mohammed Aadil and Jonathan Chacko

1. Core Machine Learning Model(s) / AI Components

- **Vertex Document AI (Pre-Trained Model):**
    - Vertex has Document AI which Provides Us with Option of Training or Selecting a Pre-Trained Model to Use.
    - Since Expense Bills are available under the Pre-Trained Model we selected it and have the Option to Fine Tune it in the Future.
- **Classification**:
    - Implemented a highly structured ADK (Agent Development Kit) model pipeline for receipt item classification.
    - Pipeline uses strict, well-defined input and output formats for reliable and explainable classification results.
    - Switched from the previous Ollama-based LLM approach (Milestone 2) to the new multi-agent ADK pipeline for better control, structure, and validation.

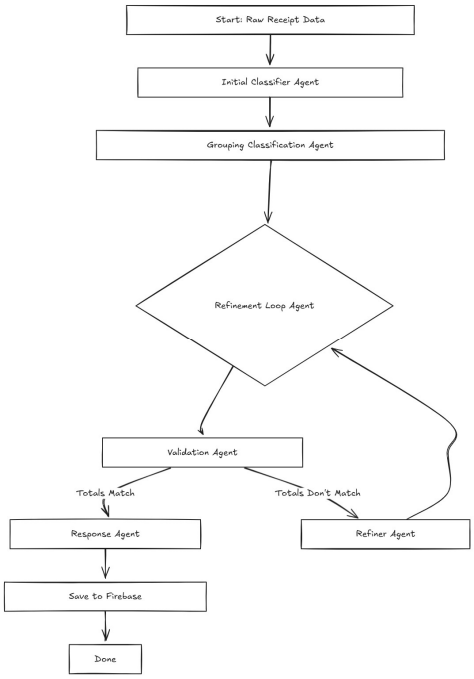

- **Regression Model (Upcoming)**:
    - A regression model for spend prediction is planned and will be integrated in the next stage (along with the Dashboard).

2. System Integration

- **Integration Demo**:
    - The gcp_adk_classification.py module demonstrates seamless integration between the main application and the ADK classification pipeline.
    - End-to-end flow: Discord bot → Flask API (main_api.py) → GCP Document AI OCR → ADK agent pipeline → Firebase storage.
- **Upcoming Dashboard**:
    - A user dashboard to visualize classified and predicted data is planned for the next milestone.
- **Screenshots from the Demo of the Deployed system so far:**

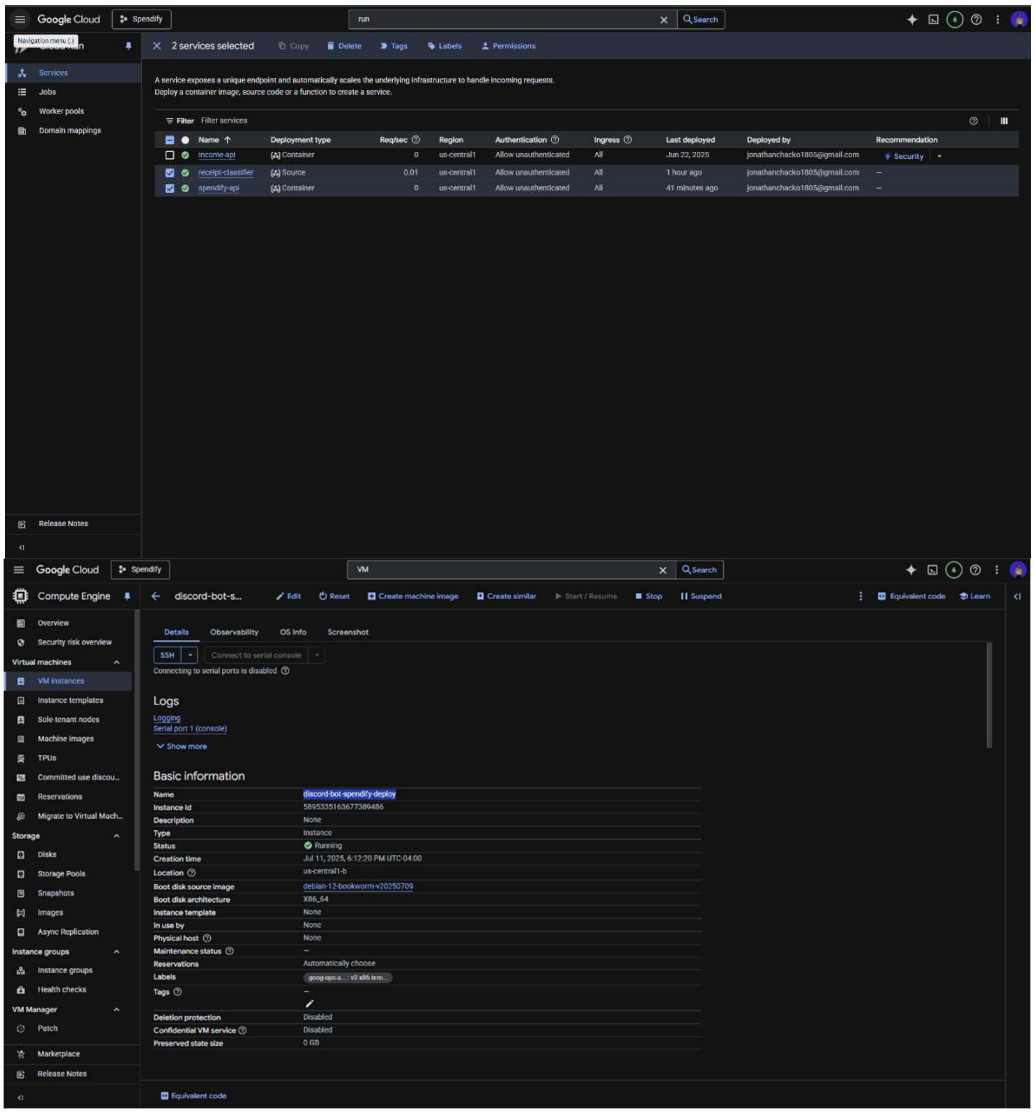

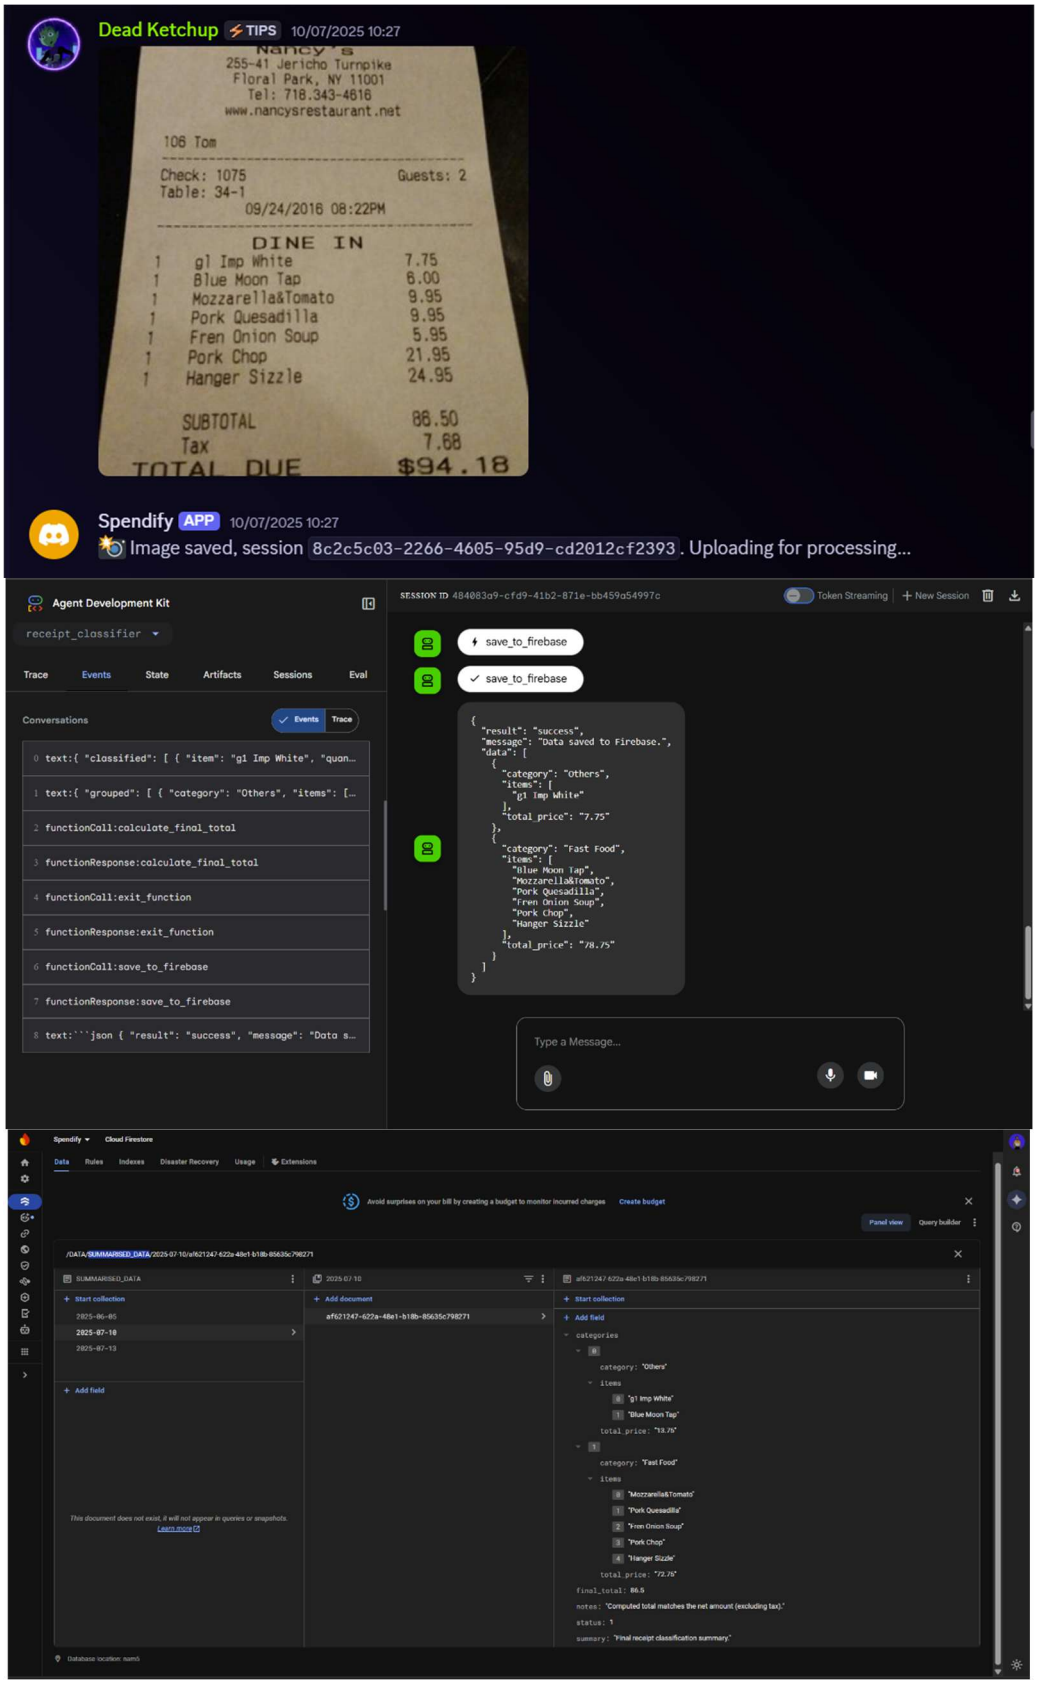
Dead Ketchup ⚡TIPS 10/07/2025 10:27
255-41 Jericho Turnpike
Floral Park, NY 11001
Tel: 718.343-4616
www.nancysrestaurant.net
106 Tom
Check: 1075 Guests: 2
Table: 34-1
09/24/2016 08:22PM
DINE IN
1 gl Imp White 7.75
1 Blue Moon Tap 6.00
1 Mozzarella&Tomato 9.95
1 Pork Quesadilla 9.95
1 Fren Onion Soup 5.95
1 Pork Chop 21.95
1 Hanger Sizzle 24.95
SUBTOTAL 86.50
Tax 7.68
TOTAL DUE $94.18
Spendify APP 10/07/2025 10:27
Image saved, session 8c2c5c03-2266-4605-95d9-cd2012cf2393. Uploading for processing...
Agent Development Kit
receipt_classifier
Trace Events State Artifacts Sessions Eval
Conversations
Events Trace
0 text:{ "classified": [{ "item": "g1 Imp White", "quan...
1 text:{ "grouped": [{ "category": "Others", "items": [...
2 functionCall:calculate_final_total
3 functionResponse:calculate_final_total
4 functionCall:exit_function
5 functionResponse:exit_function
6 functionCall:save_to_firebase
7 functionResponse:save_to_firebase
8 text:```json { "result": "success", "message": "Data s...
SESSION ID 484083a9-cfd9-41b2-871e-bb459a54997c
Token Streaming
New Session
save_to_firebase
save_to_firebase
{
"result": "success",
"message": "Data saved to Firebase.",
"data": [
{
"category": "Others",
"items": [
"g1 Imp White"
],
"total_price": "7.75"
},
{
"category": "Fast Food",
"items": [
"Blue Moon Tap",
"Mozzarella&Tomato",
"Pork Quesadilla",
"Fren Onion Soup",
"Pork Chop",
"Hanger Sizzle"
],
"total_price": "78.75"
}
]
}
Type a Message...
Spendify
Cloud Firestore
Data Rules Indexes Disaster Recovery Usage Extensions
Avoid surprises on your bill by creating a budget to monitor incurred charges
Create budget
Panel view
Query builder
/DATA/SUMMARISED_DATA/2025-07-10/af621247-622a-48e1-b18b-85635c798271
SUMMARISED_DATA
2025-07-10
af621247-622a-48e1-b18b-85635c798271
Start collection
Add document
Add field
2025-06-05
2025-07-10
2025-07-13
This document does not exist, it will not appear in queries or snapshots.
Learn more
categories
0
category: "Others"
items
0 "g1 Imp White"
1 "Blue Moon Tap"
total_price: "7.75"
1
category: "Fast Food"
items
0 "Mozzarella&Tomato"
1 "Pork Quesadilla"
2 "Fren Onion Soup"
3 "Pork Chop"
4 "Hanger Sizzle"
total_price: "72.75"
final_total: 86.5
notes: "Computed total matches the net amount (excluding tax)."
status: 1
summary: "Final receipt classification summary."
Database location: nam5

- The full pipeline was tested end-to-end:
    - Discord Bot uploads a receipt image
    - API receives and processes the image, runs OCR
    - Data sent to ADK classification pipeline (Initial Classifier → Grouping Agent → Validation/Refinement Loop → Final Response Agent)
    - Classification output is validated and stored in Firebase
- **Result:** API endpoint returns predictions successfully.

3. Initial Deployment Steps

Deployment guides and step-by-step instructions are provided in the following repository links. For each major component, use the commands below for deployment:

a. ADK Classification Agent

- Guide: deploy-adk.md
- **Cloud Run Deployment Command:**

```
adk deploy cloud_run \
  --project=$GOOGLE_CLOUD_PROJECT \
  --region=$GOOGLE_CLOUD_LOCATION \
  --service_name=$SERVICE_NAME \
  --app_name=$SERVICE_NAME \
  --port=8080 \
  --log_level=info \
  --with_ui \
  ./receipt_classifier
```

b. Flask API

- Guide: deploy-api.md
- **Build & Deploy to Cloud Run:**

```
gcloud builds submit --tag gcr.io/$GOOGLE_CLOUD_PROJECT/$SERVICE_NAME .
gcloud run deploy $SERVICE_NAME \
  --image gcr.io/$GOOGLE_CLOUD_PROJECT/$SERVICE_NAME \
  --platform managed \
  --region $GOOGLE_CLOUD_LOCATION \
  --allow-unauthenticated
```

c. **Discord Bot**

- Guide: deploy-bot.md
- **VM Deployment Command:**

```
ssh -i ~/.ssh/gcp_key user@VM_IP
tmux new -s discordbot
python3 bot.py
```

Each document contains Dockerfile usage, environment variable setup, GCP Cloud Run deployment, and troubleshooting tips specific to each core component.

4. Progress, Planning & Next Steps

From Milestone 2 → Milestone 3

- Major change: Replaced Ollama/LLM-based classification with the robust ADK model pipeline.
- Improved reliability, structure, and validation by enforcing strict input/output schemas and a multi-agent review/refine process.

Detailed Task Breakdown (Next 3 Weeks, by Role)

- **Regression model implementation & integration (Aliyyah):**
  - Predict user spend (time series/regression)
  - Integrate into existing backend pipeline
- **Dashboard development (Aadil):**
  - Build initial UI for receipt and prediction visualization
  - Connect dashboard to Firebase/API
- **Deployment & system integration (Jonathan):**
  - Deployment of Regression Model and Dashboard
  - Full system end-to-end testing

5. Challenges Faced to Achieve Milestone 3

- Ollama proved inaccurate and unreliable for classification, which led us to adopt GCP ADK as the primary classification engine.
- Swapping from Ollama to ADK allows us to access the larger computational power of GCP and benefit from a more scalable, production-grade environment.

- Integrating the highly structured ADK model pipeline with strict input/output, which required reworking the entire classification workflow.
- GCP ADK is an amazing tool for building agent pipelines, but it lacks comprehensive documentation and real-world examples, which significantly slowed down development.
- Deployment was difficult for GCP ADK for the same reasons as above, with a lot of trial and error required to achieve a working cloud setup.
- Ensuring compatibility and smooth data handoff between Discord Bot, Flask API, GCP Document AI, and ADK pipeline.
- Debugging Firestore data serialization and managing Firestore schema changes.
- Handling OCR inconsistencies in receipt formats from GCP Document AI.
- Adapting deployment strategies (switching from Ollama to ADK), requiring updates to containerization and cloud setup.

6. Remaining Challenges

- **Regression Model:** Need to design and validate regression predictions; may need additional user data.
- **Dashboard:** Ensuring robust connection between dashboard and backend, and handling real-time updates.
- **Validation:** Ensuring ground-truth data for accurate model validation (classification & regression).
- **Performance:** Scaling API/bot for multiple concurrent users; latency testing on cloud.